

Winter Midterm Progress Report: Multi-Camera, SoM Based, Real-Time Video Processing for UAS and VR/AR Applications

Area 51: Shu-Ping Chien, Brock Smedley, W Keith Striby Jr

06 May 2018

CS463, Senior Software Engineering Project, Spring 2018



Abstract

This document highlights the group's progress made during the first half of the Spring 2018 term. The project is introduced along with an explanation of the hardware and software surrounding our product's planned solution. Then details of the project's status, work remaining for the team, and problems experienced are elaborated on.

CONTENTS

1	Introduction	2
1.1	Purpose	2
1.2	Scope	2
1.3	Overview	2
2	Current Project Status	3
2.1	Image Processing	3
2.2	Testing	4
2.3	Troubleshooting Latency	4
2.3.1	GStreamer research	4
2.3.2	TX2 Rebuild	5
3	Work Remaining for Project	5
3.1	Troubleshooting Latency	5
3.2	Finish Latency and Frame Rate Tests	6
4	Problems Experienced and Solutions if Applicable	6
4.1	Latency in our Image Processing	6
4.2	TX2 Rebuild	8
5	Conclusion	8
6	Code Involved in the Project	8
7	Images of our Software Output	9

1 INTRODUCTION

1.1 Purpose

Our project is required to provide a video output at near real-time from a multi-camera input utilizing an NVIDIA Jetson TX1 or TX2 system-on-module (SoM). The software produced will perform image processing and edge computing on the camera input to display a visually enhanced and stitched video output. Size, weight, power, and cost (SWaP-C) requirements for the project are due to its application being for UAS and VR/AR, and from our system utilizing a mass-produced SoM, particularly the Jetson TX2.

1.2 Scope

Application specific hardware and software are vital for the project due the system requirements and constraints of the project. The NVIDIA Jetson TX2 runs on an Ubuntu based Linux for Tegra (L4T) operating system created specifically for producing customized imaging, installed on a GPU-accelerated dual-core CPU with dual Image Signal Processors (ISPs).

L4T, being an Ubuntu variant, provides a friendly development interface with access to a large repository of additional software. NVIDIA provides the Jetson Software Development Pack (JetPack) 3.1 which has L4T 28.1 and is capable of supporting a multitude of multimedia and image processing application program interfaces (APIs). Jetpack is installed on an Ubuntu host computer and is then flashed to the TX2's memory, and selected libraries are then installed when the TX2 is self-sustaining.

The cameras will utilize the most widely used camera interface for mobile applications, MIPI CSI-2, which is capable of supporting 1080p, 4k, and 8k video. A carrier board will connect the cameras to the Jetson TX2 that also provides additional application-specific interfaces and peripherals.

The software materials for this project will support the NVIDIA Jetson TX2 and be able to produce stitched images from cameras in near real-time. The data input from cameras are sent through the CSI and carrier board, and then arrive at our software in the TX2 in raw pixel form. The pixels are then processed through our GStreamer pipelines for distributing and transforming their data to the image processor. OpenCV libraries in our software is used to access the data in the image processor to produce our desired image output. Due to using CSI cameras the process avoids writing to memory and therefore reading our input data to storage which reduces latency. The image processing software will utilize the images from media stream and produce a streaming video to the display device.

1.3 Overview

This document provides a recap of the progress made on our project through the first half of the Spring 2018 term. This term the group has been focused on getting rid of our issues surrounding latency when streaming video before

and after image processing. For our current project status, we discuss what we have done to test and troubleshoot latency, the capabilities of our image processing, and testing performed so far. Then the work remaining for the project is elaborated on, which focuses on more troubleshooting of our latency issues and finishing testing of our product. Then we will discuss problems experienced and solutions where applicable. Finally, the report concludes with our three-camera GStreamer pipeline code and images of our software output.

2 CURRENT PROJECT STATUS

Most base goals of the project established by the Requirements Document have been met, and are waiting to complete the rest following additional troubleshooting. So far this term the main focus has been working on reducing and eliminating issues with latency in video processing. This section elaborates on the current status of image processing, testing, and troubleshooting latency.

2.1 Image Processing

The image processing for both stitching and tiling applications executes with OpenCV, which is a C++ library intended for real-time programs. The stitching application combines two overlapping images by the built-in `stitch` function in OpenCV, and it works well by automatically combining the same pixels from two video streams. Unfortunately the speed of this application is no where near real-time, which is expected by the Requirements Document. The following code shows how the `stitch` function works.

```

imgs.push_back(fr1);
imgs.push_back(fr2);
Ptr<Stitcher> test = Stitcher::create(mode, try_use_gpu);
Stitcher::Status status = test->stitch(imgs, pano);

if (status != Stitcher::OK)
{
    cout << "Error_stitching_Code:" <<int(status)<<endl;
    return -1;
}
imshow("Stitched_Image", pano);

```

The second application tiles video streams with the built-in `hconcat` and `vconcat` functions in OpenCV. The `hconcat` function combines two and three video streams horizontally, and the `vconcat` combines them vertically. These functions worked more efficient and there was little latency added by these image processing applications. The code below shows how three camera inputs are tiled into a single window.

```
cap1 >> fr1;
cap2 >> fr2;
cap3 >> fr3;
hconcat(fr1, fr2, tile);
hconcat(tile, fr3, out);
imshow("Tiled_Image", out);
```

2.2 Testing

In order to prove the efficiency of the image processing applications, testing to record latency is required to be performed. The latency is proven by comparing the stop-watch time output on a display connected to the Jetson TX2 before and after image processing. Testing on the image processing applications have yet to be performed due to the latency that exists before the `stitch`, `hconcat`, and `vconcat` parts of the programs. The GStreamer pipelines have latency that increases when more cameras are added for input, with one camera in 0.088 sec, two cameras in 0.293 sec, and 0.52 sec for three cameras.

Due to the issues in latency frame rate testing has been added to check the speed that images in the video stream are displayed. The frame rate testing has been setup to show the timing of the cameras before and after the GStreamer pipelines. The results for all camera setups were nearly the same with frames per second being between 30 and 31.

2.3 Troubleshooting Latency

Since a specific requirement in the project was asked for the video output with near real-time performance, different angles to troubleshoot for latency has been performed. The two main topics were proposed on GStreamer pipeline research and a TX2 rebuild in order to accelerate the image capturing speed with more potential functionalities.

2.3.1 GStreamer research

The first attempt for the group was to have progression on GStreamer pipeline research and manipulation. In this project, the `nvcamerasrc` was used to capture video input in the GStreamer pipeline, so most of resources about `nvcamerasrc` and GStreamer were available from NVIDIA. Then some test about breaking and adjustment on the current pipeline was also processed through the Accelerated GStreamer User Guide from NVIDIA developer in order to collect more data and comparisons. However, there was very little found on improving latency following web researches. Help was pursued on the NVIDIA developer forum but no guidance has been given.

2.3.2 TX2 Rebuild

Due to the loosely structured nature of non-versioned project collaboration, our first TX2 system build was a bit unstable. To resolve any system configuration errors and ensure that the new build was properly configured, the team decided to procure another TX2 and rebuild it from scratch. In order to do this reliably, the team dedicated one member to build the new system and document the details in order to verify 1. That the system wasn't causing issues with our stitching software, and 2. That we could reliably rebuild a working module without any system errors.

The first part was the simplest to verify: once the new build was completed, we tested the new module against our image processing software and verified that the problem did in fact lie within our implementation, and not in the system configuration. The second part, documenting the process, was also very simple. Bash scripts can be used to automate a large portion of the build process, including launching Jetpack, flashing the module with the Spacely board support image, and installing software libraries on the newly-installed operating system. We simply copied the commands which we originally used to build the platform into an executable file, which we then could run with one simple command.

3 WORK REMAINING FOR PROJECT

In the remaining weeks that we have to work on the software of our project we hope to solve the latency in our GStreamer pipeline and image processing. Questions surrounding our existing pipelines will continued to be asked and therefore adjusted. If our troubleshooting fails or after the hopeful success of fixing it, we will complete latency and frame rate testing of our software. This section elaborates on our plan of attack to complete these tasks.

3.1 Troubleshooting Latency

Although we have scoured the internet in search of solutions to our latency issues we must continue this effort until we are no longer allowed to. Troubleshooting has isolated the issue in two areas, the first being in our GStreamer pipelines, and the second occurring in image processing. This is how our troubleshooting will continue to be isolated and attacked, and this will continue until further isolation is determined or if our software requires any major overhaul.

In GStreamer pipelines commands are noted with the exclamation mark, and these denote filters or pads. The pipelines in our software have several of these pads to mutate the input coming from the cameras so the data ready for image processing. We are continuing to tinker with and manipulate our pads to see if there is any change in system response.

Research surrounding more solutions to our GStreamer latency has not turned up much due to the lack of elaboration we are able to find. Resources on GStreamer tend to show only examples with limited explanation, making it difficult for us to determine where our issue is. Our search for answers will continue looking for resources elaborating on GStreamer

pipelines.

Latency surrounding our image processing software will be sought after when we solve the issue of our GStreamer pipelines. We are unable to determine the amount of latency that our image processing of stitched and tiled videos have and therefore would be an ineffective use of our time. Solutions have been found regarding our slow image processing but not implemented, and we will determine what path to take when our GStreamer latency issues are solved based on how much latency remains, if any.

3.2 Finish Latency and Frame Rate Tests

The completion of our latency testing has been delayed in hopes of finding a solution to the latency in our GStreamer pipeline. Since we know and understand the amount of latency that exists in the GStreamer pipeline all that is remaining is determining the amount of latency in our stitching and tiling software.

Our latency tests are fairly simple to setup and perform. With the software running and producing a video output, a running internet stop-watch is placed on the output screen in a window next to it. With the cameras pointed at the output screen showing the stop-watch we then take a picture of these two windows side-by-side. The difference in time captured in our picture shows the amount of latency between the stop-watch in real-time and following image processing.

Accompanied with our latency tests we have decided to also produce frame rate testing for our client. The frame rate of our software was questioned during the troubleshooting of our latency issues, and since frame rate is part of our GStreamer pipeline it is easily adjusted and displayed on the output display. This test has been performed with a two and three-camera setup of our GStreamer output and produces the expected result of 30 frames per second.

4 PROBLEMS EXPERIENCED AND SOLUTIONS IF APPLICABLE

Over the duration of the spring term the team has been battling with fine tuning and improving software created for image processing. The initial versions of our software to produce software that stitches and tiles video output were not what our client expected of our product. The team has had trouble improving the software so that the output is near real-time. This section details the various problems we have encountered during the course of development during the term thus far, and our attempts and ideas about solutions to these problems.

4.1 Latency in our Image Processing

Our image processing software developed to stitch or combine two overlapping images into one wider image, suffers from high latency and low frame rate. The team has identified three potential culprits that are elaborated on below: the

GStreamer pipeline, a GStreamer plugin `nvcamerasrc`, and the OpenCV stitcher class and function.

GStreamer: It is possible that GStreamer is incapable of processing the camera data at the bandwidth required to produce a fair-quality image (1280x720) at a reasonable frame rate (≥ 30 fps). Although GStreamer is a widely-used image processing software and should be able to handle the bandwidth, we have not had the chance to look into alternatives until recently. OpenCV allows for the use of Video4Linux, but since we only had one working build we decided not to test an alternate configuration for fear of ruining our only working version. We also have the option of using libargus instead of GStreamer, and there is also the option to run a native V4L2 (Video4Linux 2) application (which appears to offer the lowest latency by way of structure; see Figure 1).

nvcamerasrc: The GStreamer plugin `nvcamerasrc` is designed to deliver the video stream to our image processing software, and may not be suitable for high-bandwidth processing. The `nvcamerasrc` plugin offers Image Signal Processor control, which allows us to adjust cameras for lighting and color conditions. However, for our product this may not be useful and other options available should be pursued.

In Figure 1 we see alternatives, the `v4l2src` GStreamer plugin. The `v4l2src` plugin connects directly to the V4L Mediacontroller framework bypassing the Camera Core software layer and libargus, which is completely separate from GStreamer.

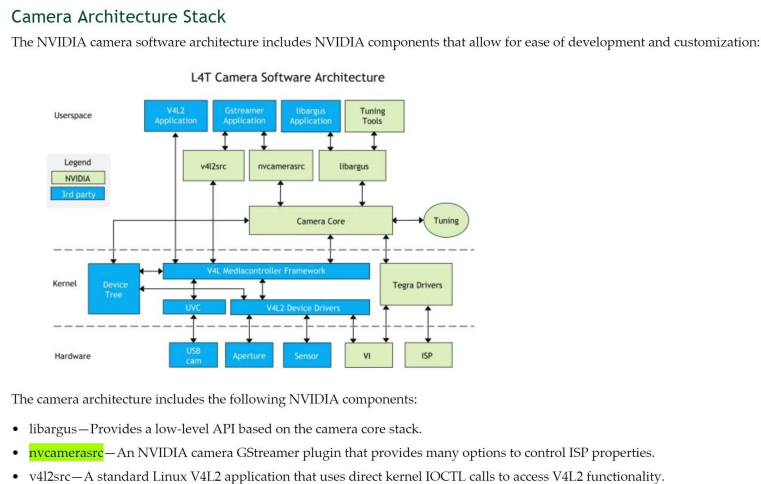


Fig. 1: TX2 Camera Interface Architecture

OpenCV Stitcher: It is possible that our implementation of the OpenCV Stitcher class and function may be less-than-optimal. While the group did program it to use the GPU which is expected to be the fastest option, there are still a lot of moving parts that could be inhibiting the system's performance. In the official OpenCV 3.3 documentation, there is a "mode" option that specifies whether to use a 'panorama' or a 'scan.' According to the documentation, panoramas, which we currently use, are best used for photographic panoramas. The 'scan' option uses an Affine Transformation algorithm, which may offer better performance. OpenCV Stitcher also allows for fine-tuned control, as described in the documentation.

4.2 TX2 Rebuild

Rebuilding the TX2 was an arduous task. Our original system, as mentioned previously, was built collaboratively, with all changes being recorded into a Google Doc. In order to build a new, reliable system, we had to pore through our progress document (which grew to 30 pages) and cipher out which bits were needed, and which were not. Unfortunately, some software conflicts arose between our configuration and the configuration described in various bits of official documentation, and the system had to be slowly, gradually built up. In the interest of saving time, the group recorded all working configuration changes; software [dependency] installs, environment variable settings, and the order which they should run in; into bash scripts, which ensured that the configuration could be repeatedly rebuilt without any discrepancies in the process, which are caused by human error.

5 CONCLUSION

With the latest progression with the project, the group currently had a compatible carrier board and CSI-2 cameras with the TX2, and which was able to achieve multi-camera image processing. The software stitched a video output from two cameras and produced a two and three-camera tiled video output. Also, the latency and frame rate has been estimated for image capturing through GStreamer pipeline.

During the sprint term, the group had most attempts on improving the latency problem and provided several potential solutions. Although the current result on the project was still hard to achieve near real-time, the group has had some progression on latency. Once the final test would be done with the latest program, the group then calculated the latency for stitched and tiled video output.

6 CODE INVOLVED IN THE PROJECT

Below is our GStreamer pipeline that processes a three-camera input and displays their feed to an output screen. Each camera input starts with `nvcamerasrc sensor-id=` and then is manipulated and filtered through their respective series of pads.

Listing 1: Three-Camera GStreamer Pipeline

```
DISPLAY=:0 gst-launch-1.0 nvcamerasrc sensor-id=0 fpsRange="30_30" !
'video/x-raw(memory:NMM),_width=(int)640,_height=(int)480,_format=(string)I420 ,
framerate=(fraction)30/1' ! nvegltransform ! nveglglessink nvcamerasrc sensor-id=2
fpsRange="30_30" ! 'video/x-raw(memory:NMM),_width=(int)640,_height=(int)480,
format=(string)I420,_framerate=(fraction)30/1' ! nvegltransform ! nveglglessink
nvcamerasrc sensor-id=1 fpsRange="30_30" ! 'video/x-raw(memory:NMM),_width=(int)640,
height=(int)480,_format=(string)I420 ,_framerate=(fraction)30/1' ! nvegltransform !
nveglglessink -e
```

7 IMAGES OF OUR SOFTWARE OUTPUT




Fig. 2: Video output of three cameras using GStreamer.



Fig. 3: Video output of our stitching program using OpenCV.

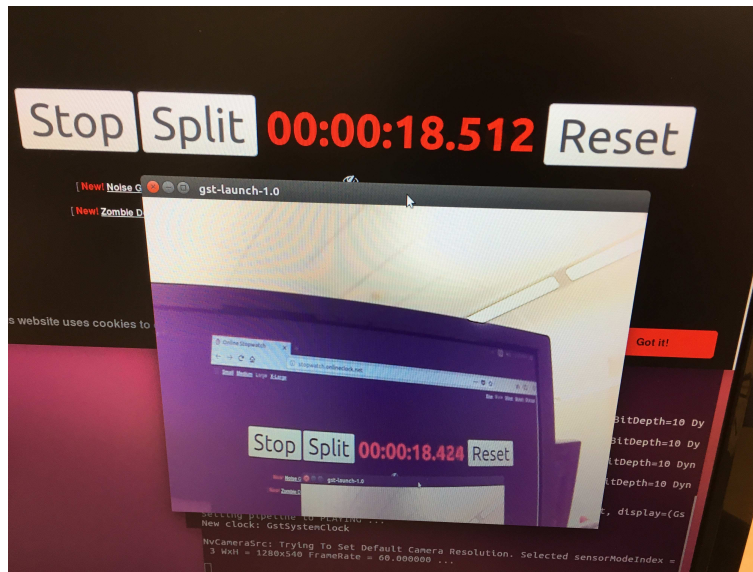


Fig. 4: The result of our one-camera latency test.

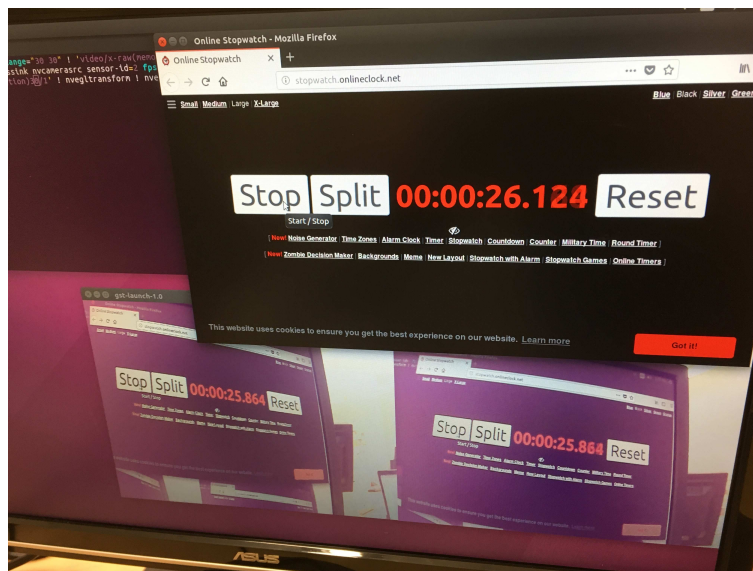


Fig. 5: The result of our two-camera latency test.




Fig. 6: The result of our three-camera latency test.




Fig. 7: The result of our frame per second test.