

# Winter Midterm Progress Report: Multi-Camera, SoM Based, Real-Time Video Processing for UAS and VR/AR Applications

Area 51: Shu-Ping Chien, Brock Smedley, W Keith Striby Jr

06 May 2018

CS463, Senior Software Engineering Project, Spring 2018



## **Abstract**

This document highlights the group's progress made during the first half of the Spring 2018 term. The project is introduced along with an explanation of the hardware and software surrounding our product's planned solution.

### **EDIT THIS:**

Then each group member explains their contributions to the project, what they have remaining to complete the projects requirements, and problems experienced with solutions if applicable.

## CONTENTS

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Purpose . . . . .	2
1.2	Scope . . . . .	2
1.3	Overview . . . . .	3
<b>2</b>	<b>Current Project Status</b>	<b>3</b>
2.1	Image Processing . . . . .	3
2.2	Testing . . . . .	3
2.3	Troubleshooting Latency . . . . .	3
2.3.1	GStreamer research . . . . .	3
2.3.2	TX2 Rebuild . . . . .	3
<b>3</b>	<b>Work Remaining for Project</b>	<b>4</b>
3.1	Troubleshooting Latency . . . . .	4
3.2	Finish Latency and Frame Rate Tests . . . . .	5
<b>4</b>	<b>Problems Experienced and Solutions if Applicable</b>	<b>5</b>
4.1	Latency in our Image Processing . . . . .	5
4.2	TX2 Rebuild . . . . .	6
<b>5</b>	<b>Conclusion</b>	<b>6</b>
<b>6</b>	<b>Code Involved in the Project</b>	<b>6</b>
<b>7</b>	<b>Images of our Software Output</b>	<b>7</b>

# 1 INTRODUCTION

## 1.1 Purpose

Our project is required to provide a video output at near real-time from a multi-camera input utilizing an NVIDIA Jetson TX1 or TX2 system-on-module (SoM). The software produced will perform image processing and edge computing on the camera input to display a visually enhanced and stitched video output. Size, weight, power, and cost (SWaP-C) requirements for the project are due to its application being for UAS and VR/AR, and from our system utilizing a mass-produced SoM, particularly the Jetson TX2.

## 1.2 Scope

Application specific hardware and software are vital for the project due to the system requirements and constraints of the project. The NVIDIA Jetson TX2 runs on an Ubuntu based Linux for Tegra (L4T) operating system created specifically for producing customized imaging, installed on a GPU-accelerated dual-core CPU with dual Image Signal Processors (ISPs).

L4T, being an Ubuntu variant, provides a friendly development interface with access to a large repository of additional software. NVIDIA provides the Jetson Software Development Pack (JetPack) 3.1 which has L4T 28.1 and is capable of supporting a multitude of multimedia and image processing application program interfaces (APIs). Jetpack is installed on an Ubuntu host computer and is then flashed to the TX2's memory, and selected libraries are then installed when the TX2 is self-sustaining.

The cameras will utilize the most widely used camera interface for mobile applications, MIPI CSI-2, which is capable of supporting 1080p, 4k, and 8k video. A carrier board will connect the cameras to the Jetson TX2 that also provides additional application-specific interfaces and peripherals.

The software materials for this project will support the NVIDIA Jetson TX2 and be able to produce stitched images from cameras in near real-time. The data input from cameras are sent through the CSI and carrier board, and then arrive at our software in the TX2 in raw pixel form. The pixels are then processed through our GStreamer pipelines for distributing and transforming their data to the image processor. OpenCV libraries in our software is used to access the data in the image processor to produce our desired image output. Due to using CSI cameras the process avoids writing to memory and therefore reading our input data to storage which reduces latency. The image processing software will utilize the images from media stream and produce a streaming video to the display device.

### 1.3 Overview

This document provides a recap of the progress made on our project through the first half of the Spring 2018 term. This term the group has been focused on getting rid of our issues surrounding latency when streaming video before and after image processing. For our current project status we discuss what we have done to test and troubleshoot latency, the capabilities of our image processing, and testing performed so far. Then the work remaining for the project is elaborated on, which focuses on more troubleshooting of our latency issues and finishing testing of our product. Then we will discuss problems experienced and solutions where applicable. Finally the report concludes with our three-camera GStreamer pipeline code and images of our software output.

## 2 CURRENT PROJECT STATUS

intro

### 2.1 Image Processing

types

### 2.2 Testing

stuff

### 2.3 Troubleshooting Latency

stuff

#### 2.3.1 GStreamer research

stuff

#### 2.3.2 TX2 Rebuild

Due to the loosely structured nature of non-versioned project collaboration, our first TX2 system build was a bit unstable. To resolve any system configuration errors and ensure that the new build was properly configured, the team decided to procure another TX2 and rebuild it from scratch. In order to do this reliably, the team dedicated one member to build

the new system and document the details in order to verify 1. That the system wasn't causing issues with our stitching software, and 2. That we could reliably rebuild a working module without any system errors.

The first part was the simplest to verify: once the new build was completed, we tested the new module against our image processing software and verified that the problem did in fact lie within our implementation, and not in the system configuration. The second part, documenting the process, was also very simple. Bash scripts can be used to automate a large portion of the build process, including launching Jetpack, flashing the module with the Spacely board support image, and installing software libraries on the newly-installed operating system. We simply copied the commands which we originally used to build the platform into an executable file, which we then could run with one simple command.

### **3 WORK REMAINING FOR PROJECT**

In the remaining weeks that we have to work on the software of our project we hope to solve the latency in our GStreamer pipeline and image processing. Questions surrounding our existing pipelines will continued to be asked and therefore adjusted. If our troubleshooting fails or after the hopeful success of fixing it, we will complete latency and frame rate testing of our software. This section elaborates on our plan of attack to complete these tasks.

#### **3.1 Troubleshooting Latency**

Although we have scoured the internet in search of solutions to our latency issues we must continue this effort until we are no longer allowed to. Troubleshooting has isolated the issue in two areas, the first being in our GStreamer pipelines, and the second occurring in image processing. This is how our troubleshooting will continue to be isolated and attacked, and this will continue until further isolation is determined or if our software requires any major overhaul.

In GStreamer pipelines commands are noted with the exclamation mark, and these denote filters or pads. The pipelines in our software have several of these pads to mutate the input coming from the cameras so the data ready for image processing. We are continuing to tinker with and manipulate our pads to see if there is any change in system response.

Research surrounding more solutions to our GStreamer latency has not turned up much due to the lack of elaboration we are able to find. Resources on GStreamer tend to show only examples with limited explanation, making it difficult for us to determine where our issue is. Our search for answers will continue looking for resources elaborating on GStreamer pipelines.

Latency surrounding our image processing software will be sought after when we solve the issue of our GStreamer pipelines. We are unable to determine the amount of latency that our image processing of stitched and tiled videos have and therefore would be an ineffective use of our time. Solutions have been found regarding our slow image processing

but not implemented, and we will determine what path to take when our GStreamer latency issues are solved based on how much latency remains, if any.

### 3.2 Finish Latency and Frame Rate Tests

The completion of our latency testing has been delayed in hopes of finding a solution to the latency in our GStreamer pipeline. Since we know and understand the amount of latency that exists in the GStreamer pipeline all that is remaining is determining the amount of latency in our stitching and tiling software.

Our latency tests are fairly simple to setup and perform. With the software running and producing a video output, a running internet stop-watch is placed on the output screen in a window next to it. With the cameras pointed at the output screen showing the stop-watch we then take a picture of these two windows side-by-side. The difference in time captured in our picture shows the amount of latency between the stop-watch in real-time and following image processing.

Accompanied with our latency tests we have decided to also produce frame rate testing for our client. The frame rate of our software was questioned during the troubleshooting of our latency issues, and since frame rate is part of our GStreamer pipeline it is easily adjusted and displayed on the output display. This test has been performed with a two and three-camera setup of our GStreamer output and produces the expected result of 30 frames per second.

## 4 PROBLEMS EXPERIENCED AND SOLUTIONS IF APPLICABLE

Over the duration of the spring term the team has been battling with fine tuning and improving software created for image processing. The initial versions of our software to produce software that stitches and tiles video output were not what our client expected of our product. The team has had trouble improving the software so that the output is near real-time. This section details the various problems we have encountered during the course of development during the term thus far, and our attempts and ideas about solutions to these problems.

### 4.1 Latency in our Image Processing

Our image processing software developed to stitch or combine two overlapping images into one wider image, suffers from high latency and low frame rate. The team has identified three potential culprits that are elaborated on below: the GStreamer pipeline, a GStreamer plugin `nvcamerasrc`, and the OpenCV stitcher class and function.

GStreamer: It is possible that GStreamer is incapable of processing the camera data at the bandwidth required to produce a fair-quality image (1280x720) at a reasonable frame rate ( $\geq 30$  fps). Although GStreamer is a widely-used image processing software and should be able to handle the bandwidth, we have not had the chance to look into alternatives until recently. OpenCV allows for the use of Video4Linux, but since we only had one working build we

decided not to test an alternate configuration for fear of ruining our only working version. We also have the option of using libargus instead of GStreamer, and there is also the option to run a native V4L2 (Video4Linux 2) application (which appears to offer the lowest latency by way of structure; see Figure 1).

`nvcamerasrc`: The GStreamer plugin `nvcamerasrc` is designed to deliver the video stream to our image processing software, and may not be suitable for high-bandwidth processing. The `nvcamerasrc` plugin offers Image Signal Processor control, which allows us to adjust cameras for lighting and color conditions. However, for our product this may not be useful and other options available should be pursued.

In Figure 1 we see an alternatives, the `v4l2src` GStreamer plugin. The `v4l2src` plugin connects directly to the V4L Mediacontroller framework bypassing the Camera Core software layer and libargus, which is completely separate from GStreamer.

OpenCV Stitcher: It is possible that our implementation of the OpenCV Stitcher class and function may be less-than-optimal. While the group did program it to use the GPU which is expected to be the fastest option, there are still a lot of moving parts that could be inhibiting the system's performance. In the official OpenCV 3.3 documentation, there is a "mode" option that specifies whether to use a 'panorama' or a 'scan.' According to the documentation, panoramas, which we currently use, are best used for photographic panoramas. The 'scan' option uses an Affine Transformation algorithm, which may offer better performance. OpenCV Stitcher also allows for fine-tuned control, as described in the documentation.

## 4.2 TX2 Rebuild

Rebuilding the TX2 was an arduous task. Our original system, as mentioned previously, was built collaboratively, with all changes being recorded into a Google Doc. In order to build a new, reliable system, we had to pore through our progress document (which grew to 30 pages) and cipher out which bits were needed and which were not. Unfortunately, some software conflicts arose between our configuration and the configuration described in various bits of official documentation, and the system had to be slowly, gradually built up. In the interest of saving time, the group recorded all working configuration changes; software [dependency] installs, environment variable settings, and the order which they should run in; into bash scripts, which ensured that the configuration could be repeatedly rebuilt without any discrepancies in the process, which are caused by human error.

## 5 CONCLUSION

more words here

## 6 CODE INVOLVED IN THE PROJECT

Below is our GStreamer pipeline that processes a three-camera input and displays their feed to an output screen. Each camera input starts with `nvcamerasrc sensor-id=` and then is manipulated and filtered through their respective

series of pads.

Listing 1: Three-Camera GStreamer Pipeline

```
DISPLAY=:0 gst-launch-1.0 nvcamerasrc sensor-id=0 fpsRange="30_30" !
'video/x-raw(memory:MMU),width=(int)640,height=(int)480,format=(string)I420,
framerate=(fraction)30/1' ! nvegltransform ! nveglglessink nvcamerasrc sensor-id=2
fpsRange="30_30" ! 'video/x-raw(memory:MMU),width=(int)640,height=(int)480,
format=(string)I420,framerate=(fraction)30/1' ! nvegltransform ! nveglglessink
nvcamerasrc sensor-id=1 fpsRange="30_30" ! 'video/x-raw(memory:MMU),width=(int)640,
height=(int)480,format=(string)I420,framerate=(fraction)30/1' ! nvegltransform !
nveglglessink -e
```

## 7 IMAGES OF OUR SOFTWARE OUTPUT




Fig. 1: Video output of three cameras using GStreamer.






Fig. 2: Video output of our stitching program using OpenCV.

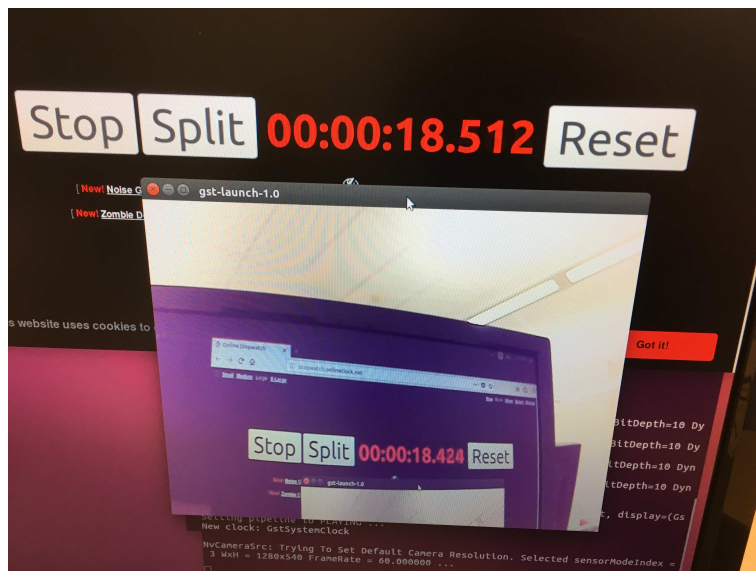


Fig. 3: The result of our one-camera latency test.

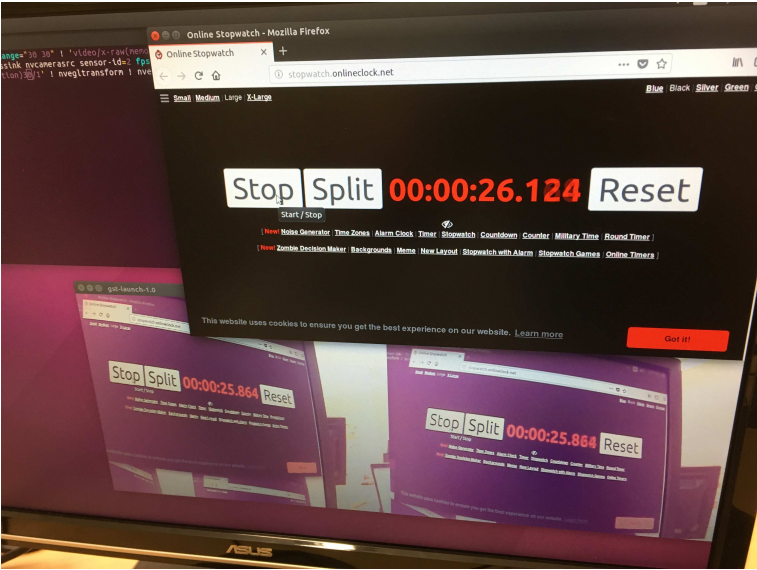


Fig. 4: The result of our two-camera latency test.




Fig. 5: The result of our three-camera latency test.




Fig. 6: The result of our frame per second test.